

5. Choose an appropriate application, find a solution using linear regression and express its performance measures.

Car Price Prediction

```
# importing necessary packages

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
%matplotlib inline
```

```
cars_info = pd.read_csv("/content/drive/MyDrive/AI-ML/DM/CarPrice_Assignment.csv")
cars_info.head()
```

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	carbody	d
0	1	3	alfa-romero giulia	gas	std	two	convertible	
1	2	3	alfa-romero stelvio	gas	std	two	convertible	
2	3	1	alfa-romero Quadrifoglio	gas	std	two	hatchback	
3	4	2	audi 100 ls	gas	std	four	sedan	
4	5	2	audi 100ls	gas	std	four	sedan	

```
# basic infos
```

```
cars_info.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   car_ID                205 non-null   int64
1   symboling              205 non-null   int64
2   CarName                205 non-null   object
3   fueltype               205 non-null   object
4   aspiration              205 non-null   object
5   doornumber             205 non-null   object
6   carbody                205 non-null   object
7   drivewheel             205 non-null   object
```

```

8  enginelocation  205 non-null  object
9  wheelbase      205 non-null  float64
10 carlength      205 non-null  float64
11 carwidth       205 non-null  float64
12 carheight      205 non-null  float64
13 curbweight     205 non-null  int64
14 enginetype     205 non-null  object
15 cylindernumber  205 non-null  object
16 enginesize      205 non-null  int64
17 fuelsystem     205 non-null  object
18 boreratio      205 non-null  float64
19 stroke         205 non-null  float64
20 compressionratio 205 non-null  float64
21 horsepower     205 non-null  int64
22 peakrpm        205 non-null  int64
23 citympg        205 non-null  int64
24 highwaympg     205 non-null  int64
25 price          205 non-null  float64
dtypes: float64(8), int64(8), object(10)
memory usage: 41.8+ KB

```

```
cars_info.describe()
```

	car_ID	symboling	wheelbase	carlength	carwidth	carheight	curbwe
count	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000
mean	103.000000	0.834146	98.756585	174.049268	65.907805	53.724878	2555.54
std	59.322565	1.245307	6.021776	12.337289	2.145204	2.443522	520.68
min	1.000000	-2.000000	86.600000	141.100000	60.300000	47.800000	1488.00
25%	52.000000	0.000000	94.500000	166.300000	64.100000	52.000000	2145.00
50%	103.000000	1.000000	97.000000	173.200000	65.500000	54.100000	2414.00
75%	154.000000	2.000000	102.400000	183.100000	66.900000	55.500000	2935.00
max	205.000000	3.000000	120.900000	208.100000	72.300000	59.800000	4066.00

```
# checking for null values
```

```
cars_info.isnull().sum()
```

```

car_ID          0
symboling       0
CarName         0
fueltype        0
aspiration      0
doornumber      0
carbody         0
drivewheel      0
enginelocation  0
wheelbase       0
carlength       0
carwidth        0
carheight       0

```

```

curbweight      0
enginetype      0
cylindernumber  0
enginesize      0
fuelsystem      0
boreratio       0
stroke         0
compressionratio 0
horsepower      0
peakrpm        0
citympg        0
highwaympg     0
price          0
dtype: int64

```

```

#Check the correlation between each of the columns
cars_info.corr()

```

```

# There is good correlation between
# 1. engine size and horse power
# 2. curb weight and car length
# 3. car width and car length
# 4. Engine size and price
# 5. curbweight and price


```

	car_ID	symboling	wheelbase	carlength	carwidth	carheight
car_ID	1.000000	-0.151621	0.129729	0.170636	0.052387	0.255960
symboling	-0.151621	1.000000	-0.531954	-0.357612	-0.232919	-0.541038
wheelbase	0.129729	-0.531954	1.000000	0.874587	0.795144	0.589435
carlength	0.170636	-0.357612	0.874587	1.000000	0.841118	0.491029
carwidth	0.052387	-0.232919	0.795144	0.841118	1.000000	0.279210
carheight	0.255960	-0.541038	0.589435	0.491029	0.279210	1.000000
curbweight	0.071962	-0.227691	0.776386	0.877728	0.867032	0.295572
enginesize	-0.033930	-0.105790	0.569329	0.683360	0.735433	0.067149
boreratio	0.260064	-0.130051	0.488750	0.606454	0.559150	0.171071
stroke	-0.160824	-0.008735	0.160959	0.129533	0.182942	-0.055307
compressionratio	0.150276	-0.178515	0.249786	0.158414	0.181129	0.261214
horsepower	-0.015006	0.070873	0.353294	0.552623	0.640732	-0.108802
peakrpm	-0.203789	0.273606	-0.360469	-0.287242	-0.220012	-0.320411
citympg	0.015940	-0.035823	-0.470414	-0.670909	-0.642704	-0.048640
highwaympg	0.011255	0.034606	-0.544082	-0.704662	-0.677218	-0.107358
price	-0.109093	-0.079978	0.577816	0.682920	0.759325	0.119336

```
# visualaizing correlation using heat map
```

```
plt.figure(figsize=(14,7))
sns.heatmap(cars_info.corr(),annot=True,cmap='coolwarm')
```


<matplotlib.axes._subplots.AxesSubplot at 0x7f90eachbd350>



```
# highest correlation is engine size vs price
```

```
# scatter plot
```

```
sns.scatterplot(x='price',y='enginesize',data=cars_info)
plt.show()
```



```
# dropping these columns
```

```
cars_info.drop(['car_ID', 'CarName', 'fuelsystem', 'enginetype'], axis=1, inplace=True)
```

```
cars_info.head()
```

	symboling	fueltype	aspiration	doornumber	carbody	drivewheel	engineloc
0	3	gas	std	two	convertible	rwd	
1	3	gas	std	two	convertible	rwd	
2	1	gas	std	two	hatchback	rwd	
3	2	gas	std	four	sedan	fwd	
4	2	gas	std	four	sedan	4wd	

```
# encoding strings to number to use it in the model
```

```
# mapping
```

```
cars_info['fueltype'] = cars_info['fueltype'].map({'gas':'0', 'diesel':'1'})
cars_info['aspiration'] = cars_info['aspiration'].map({'std':'0', 'turbo':'1'})
cars_info['doornumber'] = cars_info['doornumber'].map({'two':'2', 'four':'4'})
cars_info['carbody'] = cars_info['carbody'].map({'convertible':'0', 'hatchback':'1'})
cars_info['drivewheel'] = cars_info['drivewheel'].map({'rwd':'0', 'fwd':'1', '4wd':'2'})
cars_info['enginelocation'] = cars_info['enginelocation'].map({'front':'0', 'rear':'1'})
cars_info['cylindernumber'] = cars_info['cylindernumber'].map({'four':'4', 'six':'6'})
# checking
cars_info.head()
```

	symboling	fueltype	aspiration	doornumber	carbody	drivewheel	engineloc
0	3	0	0	2	0	0	
1	3	0	0	2	0	0	
2	1	0	0	2	1	0	
3	2	0	0	4	2	1	
4	2	0	0	4	2	2	

Model building part

```
# taking price as the label or target
```

```
from sklearn.model_selection import train_test_split
```

```
X = cars_info.drop('price', axis=1)
```

```
X = cars_info.drop('price',axis=1,
y= cars_info['price']

X_train,X_test,y_train,y_test= train_test_split(X,y,test_size=0.3,random_state=55)

# training or fitting the curve

from sklearn.linear_model import LinearRegression
lr = LinearRegression()
lr.fit(X_train,y_train, )
```

```
LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

```
# prediction part

prediction = lr.predict(X_test)
```

▼ Perfomance measures

```
# r2 score

from sklearn.metrics import r2_score
print(r2_score(y_test,prediction))

# got 78% accuracy
```

```
0.7784967344806435
```

```
# The explained_variance_score computes the explained variance regression score.

from sklearn.metrics import explained_variance_score

explained_variance_score(y_test, prediction, multioutput='raw_values')

array([0.77858187])
```

```
# The max_error function computes the maximum residual error ,
# a metric that captures the worst case error between the predicted
# value and the true value.

from sklearn.metrics import max_error
max_error(y_test, prediction)
```

```
16867.0309776942
```

```
# Mean Abosolute Error

from sklearn.metrics import mean_absolute_error
```

```
mean_absolute_error(y_test, prediction, multioutput='raw_values')
```

```
array([2343.26176044])
```

```
# Mean squared error
```

```
from sklearn.metrics import mean_squared_error
```

```
mean_squared_error(y_test, prediction)
```

```
13119405.902743611
```

```
# test data vs prediction
```

```
plt.scatter(y_test, prediction)
```

```
<matplotlib.collections.PathCollection at 0x7f90e734e610>
```

